

**Problem 1**

Loss = 2.0

Gradient with respect to  $w_3 = 12$

Gradient with respect to  $w_2 = 28$

Gradient with respect to  $w_1 = -84$

# M8-L1 Problem 1

In this problem you will solve for  $\frac{\partial L}{\partial W_2}$  and  $\frac{\partial L}{\partial W_1}$  for a neural network with two input features, a hidden layer with 3 nodes, and a single output. You will use the sigmoid activation function on the hidden layer. You are provided an input sample  $x_0$ , the current weights  $W_1$  and  $W_2$ , and the ground truth value for the sample,  $t = -2$

$$L = \frac{1}{2}e^T e$$

```
In [1]: import numpy as np

x0 = np.array([[ -2], [ -6]])

W1 = np.array([[ -2,  1], [ 3,  8], [-12,  7]])
W2 = np.array([[ -11,  2,  5]])

t = np.array([[ -2]])
```

## Define activation function and its derivative

First define functions for the sigmoid activation functions, as well as its derivative:

```
In [2]: # YOUR CODE GOES HERE
def sigmoid(x):
    return 1/(1 + np.exp(-x))

def d_sigmoid(x):
    return sigmoid(x) * (1 - sigmoid(x))
```

## Forward propagation

Using your activation function, compute the output of the network  $y$  using the sample  $x_0$  and the provided weights  $W_1$  and  $W_2$

```
In [3]: # YOUR CODE GOES HERE
hidden_input = np.dot(W1, x0)

hidden_output = sigmoid(hidden_input)

output = np.dot(W2, hidden_output)
```

## Backpropagation

Using your calculated value of  $y$ , the provided value of  $t$ , your  $\sigma$  and  $\sigma'$  function, and the provided weights  $W_1$  and  $W_2$ , compute the gradients  $\frac{\partial L}{\partial W_2}$  and  $\frac{\partial L}{\partial W_1}$ .

```
In [5]: # YOUR CODE GOES HERE
e = output - t
L = 1/2 * np.dot(e, e)

dL_dy = e
dL_dW2 = dL_dy * hidden_output.T

dL_dhidden_output = dL_dy * W2.T
dL_dhidden_input = dL_dhidden_output * d_sigmoid(hidden_input)
dL_dW1 = np.dot(dL_dhidden_input, x0.T)

print("dL_dW1: ", dL_dW1)
print("dL_dW2: ", dL_dW2)

dL_dW1:  [[ 1.59095673e+00  4.77287018e+00]
 [-9.73264513e-24 -2.91979354e-23]
 [-1.04899214e-07 -3.14697641e-07]]
dL_dW2:  [[ 8.21031503e-02  2.43316128e-24  1.04899215e-08]]
```

Processing math: 100%

# M8-L2 Problem 1

In this problem, you will create 3 regression networks with different complexities in PyTorch. By looking at the validation loss curves superimposed on the training loss curves, you should determine which model is optimal.

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import torch
from torch import nn, optim

def generate_data():
    np.random.seed(5)
    N = 25
    x = np.random.normal(np.linspace(0,1,N),0.01).reshape(-1,1)
    y = np.random.normal(np.sin(5*(x+0.082)),0.2)
    train_mask = np.zeros(N,dtype=np.bool_)
    train_mask[np.random.permutation(N)[:int(N*0.8)]] = True
    train_x, val_x = torch.Tensor(x[train_mask]), torch.Tensor(x[np.logical_not(train_mask)])
    train_y, val_y = torch.Tensor(y[train_mask]), torch.Tensor(y[np.logical_not(train_mask)])

    return train_x, val_x, train_y, val_y

def train(model, lr=0.0001, epochs=10000):
    train_x, val_x, train_y, val_y = generate_data()
    opt = optim.Adam(model.parameters(),lr=lr)
    lossfun = nn.MSELoss()
    train_hist = []
    val_hist = []

    for _ in range(epochs):
        model.train()
        loss_train = lossfun(train_y, model(train_x))
        train_hist.append(loss_train.item())

        model.eval()
        loss_val = lossfun(val_y, model(val_x))
        val_hist.append(loss_val.item())

        opt.zero_grad()
        loss_train.backward()
        opt.step()

    train_hist, val_hist = np.array(train_hist), np.array(val_hist)
    return train_hist, val_hist

def plot_loss(train_loss, val_loss):
    plt.plot(train_loss,label="Training")
    plt.plot(val_loss,label="Validation",linewidth=1)
    plt.legend()
    plt.xlabel("Epoch")
    plt.ylabel("MSE Loss")


def plot_data(model = None):
    train_x, val_x, train_y, val_y = generate_data()
    plt.scatter(train_x, train_y,s=8,label="Train Data")
    plt.scatter(val_x, val_y,s=12,marker="x",label="Validation Data",linewidths=1)

    if model is not None:
        xvals = torch.linspace(0,1,1000).reshape(-1,1)
        plt.plot(xvals.detach().numpy(),model(xvals).detach().numpy(),label="Model",color="black")

    plt.legend(loc="lower left")

def get_loss(model):
    lossfun = nn.MSELoss()
    train_x, val_x, train_y, val_y = generate_data()
    loss_train = lossfun(train_y, model(train_x))
    loss_val = lossfun(val_y, model(val_x))
    return loss_train.item(), loss_val.item()

plt.figure(figsize=(4,3),dpi=250)
plot_data()
plt.show()
```



## Coding neural networks for regression

Here, create 3 neural networks from scratch. You can use `nn.Sequential()` to simplify things. Each network should have 1 input and 1 output. After each hidden layer, apply ReLU activation. Name the models `model1`, `model2`, and `model3`, with architectures as follows:

- `model1`: 1 hidden layer with 4 neurons. That is, the network should have a linear transformation from size 1 to size 4. Then a ReLU activation should be applied. Finally, a linear transformation from size 4 to size 1 gives the network output. (Note: Your regression network should not have an activation after the last layer!)
- `model2`: Hidden sizes (16, 16). (Two hidden layers, each with 16 neurons)
- `model3`: Hidden sizes (128, 128, 128). (3 hidden layers, each with 128 neurons)

```
In [2]: # YOUR CODE GOES HERE
model1 = nn.Sequential(nn.Linear(1, 4), nn.ReLU(), nn.Linear(4, 1))

# Define model2
model2 = nn.Sequential(nn.Linear(1, 16), nn.ReLU(), nn.Linear(16, 16), nn.ReLU(), nn.Linear(16, 1))

# Define model3
model3 = nn.Sequential(nn.Linear(1, 128), nn.ReLU(), nn.Linear(128, 128), nn.ReLU(), nn.Linear(128, 128), nn.ReLU(), nn.Linear(128, 1))
```


## Training and Loss curves

The following cell calls the provided function `train` to train each of your neural network models. The training and validation curves are then displayed.

```
In [3]: hidden_layers=["(4)","(16, 16)","(128, 128, 128)"]

plt.figure(figsize=(15,3),dpi=250)
for i,model in enumerate([model1, model2, model3]):
    loss_train, loss_val = train(model)
    plt.subplot(1,3,i+1)
    plot_loss(loss_train, loss_val)
    plt.ylim(0,0.6)
    plt.title(f"Hidden layers: {hidden_layers[i]}")
```

```
plt.show()
```



## Model performance

Let's print the values of MSE on the training and testing/validation data after training. Make note of which model is "best" (has lowest testing error).


```
In [4]: for i, model in enumerate([model1, model2, model3]):  
        train_loss, val_loss = get_loss(model)  
        print(f"Model {i+1}, hidden layers {hidden_layers[i]:>15}:   Train MSE: {train_loss:.4f}   Test MSE: {val_loss:.4f}")
```

Model 1, hidden layers	(4):	Train MSE: 0.1145	Test MSE: 0.0743
Model 2, hidden layers	(16, 16):	Train MSE: 0.0058	Test MSE: 0.0251
Model 3, hidden layers	(128, 128, 128):	Train MSE: 0.0014	Test MSE: 0.0380

## Visualization

Now we can look at how good each model's predictions are. Run the following cell to generate a visualization plot, then answer the questions.

```
In [5]: plt.figure(figsize=(15,3),dpi=250)  
        for i,model in enumerate([model1, model2, model3]):  
            plt.subplot(1,3,i+1)  
            plot_data(model)  
            plt.title(f"Hidden layers: {hidden_layers[i]}")  
        plt.show()
```



## Questions

1. For the model that overfits the most, describe what happens to the loss curves while training.
  2. For the model that underfits the most, describe what happens to the loss curves while training.
  3. For the "best" model, what happens to the loss curves while training?
- 
1. model3: The training loss decreases quickly and stays very low, indicating that the model is fitting well to the training data. However, the validation loss starts to increase after an initial decrease, showing that the model is fitting to the training data too well but struggling to generalize to the validation set. This pattern is seen typically in overfitting.
  2. model1: Both training and validation losses decrease, but they do so at a relatively less pace. The losses don't reach very low values, indicating that the model fails to capture the complexity of the underlying data distribution. This shows underfitting.
  3. model2: Both the training and validation losses decrease and eventually stabilize at low values. The gap between training and validation loss remains small, indicating a good balance between fitting the training data and generalizing to unseen data. This suggests that model2 has the is designed well for the data complexity, making it the most appropriate model in this case.

## M8-L2 Problem 2

Let's revisit the material phase prediction problem once again. You will use this problem to try multi-class classification in PyTorch. You will have to write code for a classification network and for training.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
import torch
from torch import nn, optim

def plot_loss(train_loss, val_loss):
    plt.figure(figsize=(4,2),dpi=250)
    plt.plot(train_loss,label="Training")
    plt.plot(val_loss,label="Validation",linewidth=1)
    plt.legend()
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.show()

def split_data(X, Y):
    np.random.seed(100)
    N = len(Y)
    train_mask = np.zeros(N, dtype=np.bool_)
    train_mask[np.random.permutation(N)[:int(N*0.8)]] = True
    train_x, val_x = torch.Tensor(X[train_mask,:]), torch.Tensor(X[np.logical_not(train_mask),:])
    train_y, val_y = torch.Tensor(Y[train_mask]), torch.Tensor(Y[np.logical_not(train_mask)])
    return train_x, val_x, train_y, val_y

x1 = np.array([7.4881350392732475,16.351893663724194,22.427633760716436,29.04883182996897,35.03654799338904,44.4
x2 = np.array([0.11120957227224215,0.1116933996874757,0.14437480785146242,0.11818202991034835,0.0859507900573781
X = np.vstack([x1,x2]).T
y = np.array([0,2,2,2,2,2,0,2,2,2,2,0,0,2,0,1,2,0,0,1,1,1,2,0,1,0,1,1,1,0,0,1,1,1,1,])

X = torch.Tensor(X)
Y = torch.tensor(y,dtype=torch.long)

train_x, val_x, train_y, val_y = split_data(X,Y)

def plot_data(newfig=True):
    xlim = [0,52.5]
    ylim = [0,1.05]
    markers = [dict(marker="o", color="royalblue"), dict(marker="s", color="crimson"), dict(marker="D", color="
    labels = ["Solid", "Liquid", "Vapor"]

    if newfig:
        plt.figure(figsize=(6,4),dpi=250)

    x = X.detach().numpy()
    y = Y.detach().numpy().flatten()


    for i in range(1+max(y)):
        plt.scatter(x[y==i,0], x[y==i,1], s=40, **(markers[i]), edgecolor="black", linewidths=0.4,label=labels[

    plt.scatter(val_x[:,0], val_x[:,1],s=120,c="None",marker="o",edgecolors="black",label="Test point")

    plt.title("Phase of simulated material")
    plt.legend(loc="upper right")
    plt.xlim(xlim)
    plt.ylim(ylim)
    plt.xlabel("Temperature, K")
    plt.ylabel("Pressure, atm")
    plt.box(True)

def plot_model(model, res=200):
    xlim = [0,52.5]
    ylim = [0,1.05]
    xvals = np.linspace(*xlim,res)
    yvals = np.linspace(*ylim,res)
    x,y = np.meshgrid(xvals,yvals)
    XY = np.concatenate((x.reshape(-1,1),y.reshape(-1,1)),axis=1)
    XY = torch.Tensor(XY)
    color = model.predict(XY).reshape(res,res).detach().numpy()
    cmap = ListedColormap(["lightblue","lightcoral","palegreen"])
    plt.pcolor(x, y, color, shading="nearest", zorder=-1, cmap=cmap,vmin=0,vmax=2)
    return

plot_data()
plt.show()
```



## Model definition

In the cell below, complete the definition for `PhaseNet`, a classification neural network.

- The network should take in 2 inputs and return 3 outputs.
- The network size and hidden layer activations are up to you.
- Make sure to use the proper activation function (for multi-class classification) at the final layer.
- The `predict()` method has been provided, to return the integer class value. You must finish `__init__()` and `forward()`.

```
In [2]: import torch.nn.functional as F

class PhaseNet(nn.Module):
    def __init__(self):
        super().__init__()
        # YOUR CODE GOES HERE
        self.fc1 = nn.Linear(2, 16)
        self.fc2 = nn.Linear(16, 16)
        self.fc3 = nn.Linear(16, 3)

    def predict(self,X):
        Y = self(X)
        return torch.argmax(Y,dim=1)

    def forward(self,X):
        # YOUR CODE GOES HERE
        X = F.relu(self.fc1(X))
        X = F.relu(self.fc2(X))
        X = F.softmax(self.fc3(X), dim=1)

        return X
```

## Training

Most of the training code has been provided below. Please add the following where indicated:

- Define a loss function (for multiclass classification)

- Define an optimizer and call it `opt` . You may choose which optimizer.

Make sure the training curves you get are reasonable.

```
In [3]: model = PhaseNet()

lr = 0.001
epochs = 1000

# Define loss function
# YOUR CODE GOES HERE
lossfun = nn.CrossEntropyLoss()

# Define an optimizer, `opt`
# YOUR CODE GOES HERE
opt = optim.Adam(model.parameters(), lr=lr)

train_hist = []
val_hist = []

for epoch in range(epochs+1):
    model.train()
    loss_train = lossfun(model(train_x), train_y)
    train_hist.append(loss_train.item())


    model.eval()
    loss_val = lossfun(model(val_x), val_y)
    val_hist.append(loss_val.item())

    opt.zero_grad()
    loss_train.backward()
    opt.step()
    if epoch % int(epochs / 25) == 0:
        print(f"Epoch {epoch:>4} of {epochs}:   Train Loss = {loss_train.item():.4f}   Validation Loss = {loss_val.item():.4f}")

plot_loss(train_hist, val_hist)
```

Epoch	0 of 1000:	Train Loss = 1.0793	Validation Loss = 1.1506
Epoch	40 of 1000:	Train Loss = 1.0505	Validation Loss = 1.0913
Epoch	80 of 1000:	Train Loss = 1.0412	Validation Loss = 1.0853
Epoch	120 of 1000:	Train Loss = 1.0279	Validation Loss = 1.0668
Epoch	160 of 1000:	Train Loss = 1.0062	Validation Loss = 1.0368
Epoch	200 of 1000:	Train Loss = 0.9725	Validation Loss = 1.0071
Epoch	240 of 1000:	Train Loss = 0.9305	Validation Loss = 0.9750
Epoch	280 of 1000:	Train Loss = 0.8866	Validation Loss = 0.9393
Epoch	320 of 1000:	Train Loss = 0.8453	Validation Loss = 0.8974
Epoch	360 of 1000:	Train Loss = 0.8072	Validation Loss = 0.8680
Epoch	400 of 1000:	Train Loss = 0.7729	Validation Loss = 0.8476
Epoch	440 of 1000:	Train Loss = 0.7442	Validation Loss = 0.8323
Epoch	480 of 1000:	Train Loss = 0.7224	Validation Loss = 0.8224
Epoch	520 of 1000:	Train Loss = 0.7066	Validation Loss = 0.8144
Epoch	560 of 1000:	Train Loss = 0.6949	Validation Loss = 0.8105
Epoch	600 of 1000:	Train Loss = 0.6860	Validation Loss = 0.8074
Epoch	640 of 1000:	Train Loss = 0.6790	Validation Loss = 0.8054
Epoch	680 of 1000:	Train Loss = 0.6732	Validation Loss = 0.8045
Epoch	720 of 1000:	Train Loss = 0.6684	Validation Loss = 0.8042
Epoch	760 of 1000:	Train Loss = 0.6643	Validation Loss = 0.8042
Epoch	800 of 1000:	Train Loss = 0.6608	Validation Loss = 0.8044
Epoch	840 of 1000:	Train Loss = 0.6577	Validation Loss = 0.8048
Epoch	880 of 1000:	Train Loss = 0.6549	Validation Loss = 0.8055
Epoch	920 of 1000:	Train Loss = 0.6524	Validation Loss = 0.8064
Epoch	960 of 1000:	Train Loss = 0.6501	Validation Loss = 0.8076
Epoch	1000 of 1000:	Train Loss = 0.6481	Validation Loss = 0.8087






## Plot results

Plot your network predictions with the data by running the following cell. If your network has significant overfitting/underfitting, go back and retrain a new network with different layer sizes/activations.

```
In [4]: plot_data(newfig=True)
        plot_model(model)
        plt.show()
```

Phase of simulated material



# Problem 1

Consider a 2D robotic arm with 3 links. The position of its end-effector is governed by the arm lengths and joint angles as follows (as in the figure "data/robot-arm.png"):

$$x = L_1 \cos(\theta_1) + L_2 \cos(\theta_2 + \theta_1) + L_3 \cos(\theta_3 + \theta_2 + \theta_1) \quad y = L_1 \sin(\theta_1) + L_2 \sin(\theta_2 + \theta_1) + L_3 \sin(\theta_3 + \theta_2 + \theta_1)$$

In robotics settings, inverse-kinematics problems are common for setups like this. For example, suppose all 3 arm lengths are  $L_1 = L_2 = L_3 = 1$ , and we want to position the end-effector at  $(x, y) = (0.5, 0.5)$ . What set of joint angles  $(\theta_1, \theta_2, \theta_3)$  should we choose for the end-effector to reach this position?

In this problem you will train a neural network to find a function mapping from coordinates  $(x, y)$  to joint angles  $(\theta_1, \theta_2, \theta_3)$  that position the end-effector at  $(x, y)$ .

Summary of deliverables:

1. Neural network model
2. Generate training and validation data
3. Training function
4. 6 plots with training and validation loss
5. 6 prediction plots
6. Respond to the prompts

```
In [60]: import numpy as np
import matplotlib.pyplot as plt

import torch
from torch import nn, optim

class ForwardArm(nn.Module):
    def __init__(self, L1=1, L2=1, L3=1):
        super().__init__()
        self.L1 = L1
        self.L2 = L2
        self.L3 = L3
    def forward(self, angles):
        theta1 = angles[:,0]
        theta2 = angles[:,1]
        theta3 = angles[:,2]
        x = self.L1*torch.cos(theta1) + self.L2*torch.cos(theta1+theta2) + self.L3*torch.cos(theta1+theta2+theta3)
        y = self.L1*torch.sin(theta1) + self.L2*torch.sin(theta1+theta2) + self.L3*torch.sin(theta1+theta2+theta3)
        return torch.vstack([x,y]).T

def plot_predictions(model, title=""):
    fwd = ForwardArm()

    vals = np.arange(0.1, 2.0, 0.2)
    x, y = np.meshgrid(vals,vals)
    coords = torch.tensor(np.vstack([x.flatten(),y.flatten()]).T,dtype=torch.float)
    angles = model(coords)
    preds = fwd(angles).detach().numpy()

    plt.figure(figsize=[4,4],dpi=140)

    plt.scatter(x.flatten(), y.flatten(), s=60, c="None",marker="o",edgecolors="k", label="Targets")
    plt.scatter(preds[:,0], preds[:,1], s=25, c="red", marker="o", label="Predictions")
    plt.text(0.1, 2.15, f"MSE = {nn.MSELoss()(fwd(model(coords)),coords):.1e}")
    plt.xlabel("x")
    plt.ylabel("y")
    plt.xlim(-.1,2.1)
    plt.ylim(-.1,2.4)
    plt.legend()
    plt.title(title)
    plt.show()

def plot_arm(theta1, theta2, theta3, L1=1,L2=1,L3=1, show=True):
    x1 = L1*np.cos(theta1)
    y1 = L1*np.sin(theta1)
    x2 = x1 + L2*np.cos(theta1+theta2)
    y2 = y1 + L2*np.sin(theta1+theta2)
    x3 = x2 + L3*np.cos(theta1+theta2+theta3)
    y3 = y2 + L3*np.sin(theta1+theta2+theta3)
```

```

xs = np.array([0,x1,x2,x3])
ys = np.array([0,y1,y2,y3])

plt.figure(figsize=(5,5),dpi=140)
plt.plot(xs, ys, linewidth=3, markersize=5,color="gray", markerfacecolor="lightgray",marker="o",markeredgecolor="black",zorder=100)
plt.scatter(x3,y3,s=50,color="blue",marker="P",zorder=100)
plt.scatter(0,0,s=50,color="black",marker="s",zorder=-100)

plt.xlim(-1.5,3.5)
plt.ylim(-1.5,3.5)

if show:
    plt.show()

```

## End-effector position

You can use the interactive figure below to visualize the robot arm.

```

In [61]: %matplotlib inline
from ipywidgets import interact, interactive, fixed, interact_manual, Layout, FloatSlider, Dropdown

def plot_unit_arm(theta1, theta2, theta3):
    plot_arm(theta1, theta2, theta3)

slider1 = FloatSlider(value=0, min=-np.pi*0.75, max=np.pi*0.75, step=np.pi/100, description='theta1',disabled=False)
slider2 = FloatSlider(value=0, min=-np.pi*0.75, max=np.pi*0.75, step=np.pi/100, description='theta2',disabled=False)
slider3 = FloatSlider(value=0, min=-np.pi*0.75, max=np.pi*0.75, step=np.pi/100, description='theta3',disabled=False)

interactive_plot = interactive(plot_unit_arm, theta1 = slider1, theta2 = slider2, theta3 = slider3)
output = interactive_plot.children[-1]
output.layout.height = '600px'

interactive_plot

```

```

Out[61]: interactive(children=(FloatSlider(value=0.0, description='theta1', layout=Layout(width='550px'), max=2.3561944...)

```

## Neural Network for Inverse Kinematics

In this class we have mainly had regression problems with only one output. However, you can create neural networks with any number of outputs just by changing the size of the last layer. For this problem, we already know the function to go from joint angles (3) to end-effector coordinates (2). This is provided in neural network format as `ForwardArm()`.

If you provide an instance of `ForwardArm()` with an  $N \times 3$  tensor of joint angles, and it will return an  $N \times 2$  tensor of coordinates.

Here, you should create a neural network with 2 inputs and 3 outputs that, once trained, can output the joint angles (in radians) necessary to reach the input x-y coordinates.

In the cell below, complete the definition for `InverseArm()`:

- The initialization argument `hidden_layer_sizes` dictates the number of neurons per hidden layer in the network. For example, `hidden_layer_sizes=[12,24]` should create a network with 2 inputs, 12 neurons in the first hidden layer, 24 neurons in the second hidden layer, and 3 outputs.
- Use a ReLU activation at the end of each hidden layer.
- The initialization argument `max_angle` refers to the maximum bend angle of the joint. If `max_angle=None`, there should be no activation at the last layer. However, if `max_angle=1` (for example), then the output joint angles should be restricted to the interval  $[-1, 1]$  (radians). You can clamp values with the tanh function (and then scale them) to achieve this.

```

In [62]: class InverseArm(nn.Module):
    def __init__(self, hidden_layer_sizes=[24,24], max_angle = None):
        super().__init__()
        # YOUR CODE GOES HERE
        self.max_angle = max_angle

        self.layers = nn.ModuleList()

        for i in range(len(hidden_layer_sizes)):
            self.layers.append(nn.Linear(hidden_layer_sizes[i - 1] if i > 0 else 2, hidden_layer_sizes[i]))
        self.layers.append(nn.Linear(hidden_layer_sizes[-1], 3))

    def forward(self, xy):
        # YOUR CODE GOES HERE
        angles = xy

        for layer in self.layers[:-1]:
            angles = torch.relu(layer(angles))
        angles = self.layers[-1](angles)

```

```

if self.max_angle is not None:
    angles = torch.tanh(angles) * self.max_angle

return angles

```

## Generate Data

In the cell below, generate a dataset of x-y coordinates. You should use a  $100 \times 100$  meshgrid, for x and y each on the interval  $[0, 2]$ .

Randomly split your data so that 80% of points are in `X_train`, while the remaining 20% are in `X_val`. (Each of these should have 2 columns -- x and y)

```

In [63]: # YOUR CODE GOES HERE
from sklearn.model_selection import train_test_split

x = np.linspace(0, 2, 100)
y = np.linspace(0, 2, 100)
x_grid, y_grid = np.meshgrid(x, y)
coordinates = torch.tensor(np.vstack([x_grid.ravel(), y_grid.ravel()]).T, dtype = torch.float)

X_train, X_val = train_test_split(coordinates, test_size=0.2, random_state=42)

```

## Training function

Write a function `train()` below with the following specifications:

*Inputs:*

- `model`: `InverseArm` model to train
- `X_train`:  $N \times 2$  vector of training x-y coordinates
- `X_val`:  $N \times 2$  vector of validation x-y coordinates
- `lr`: Learning rate for Adam optimizer
- `epochs`: Total epoch count
- `gamma`: ExponentialLR decay rate
- `create_plot`: ( `True` / `False` ) Whether to display a plot with training and validation loss curves

*Loss function:*

The loss function you use should be based on whether the end-effector moves to the correct location. It should be the MSE between the target coordinate tensor and the coordinates that the predicted joint angles produce. In other words, if your inverse kinematics model is `model`, and `fwd` is an instance of `ForwardArm()`, then you want the MSE between input coordinates `X` and `fwd(model(X))`.

```

In [68]: from torch.optim.lr_scheduler import ExponentialLR

def train(model, X_train, X_val, lr = 0.01, epochs = 1000, gamma = 1, create_plot = True):
    # YOUR CODE GOES HERE
    X_train = torch.tensor(X_train, dtype=torch.float32)
    X_val = torch.tensor(X_val, dtype=torch.float32)

    fwd = ForwardArm()
    lossfun = nn.MSELoss()
    opt = optim.Adam(model.parameters(), lr = lr)
    scheduler = optim.lr_scheduler.StepLR(opt, step_size = 100, gamma = gamma)

    train_hist = []
    val_hist = []

    for epoch in range(epochs):
        model.train()

        opt.zero_grad()

        angles_pred_train = model(X_train)
        coords_pred_train = fwd(angles_pred_train)

        loss_train = lossfun(coords_pred_train, X_train)
        loss_train.backward()

        opt.step()

        train_hist.append(loss_train.item())

    model.eval()
    with torch.no_grad():
        predicted_angles_val = model(X_val)
        predicted_coords_val = fwd(predicted_angles_val)

```

```

        loss_val = lossfun(predicted_coords_val, X_val)
        val_hist.append(loss_val.item())

    scheduler.step()

    if epoch % (epochs // 25) == 0:
        print(f"Epoch {epoch:>4}/{epochs}: Train Loss = {loss_train.item():.4f}, Validation Loss = {loss_val.item():.4f}")

    if create_plot:
        plt.figure(figsize=(10, 5))
        plt.plot(train_hist, label='Training Loss')
        plt.plot(val_hist, label='Validation Loss')
        plt.xlabel('Epoch')
        plt.ylabel('MSE Loss')
        plt.title('Training and Validation Loss')
        plt.legend()
        plt.show()

    return model

```

## Training a model

Create 3 models of different complexities (with `max_angle=None`):

- `hidden_layer_sizes=[12]`
- `hidden_layer_sizes=[24,24]`
- `hidden_layer_sizes=[48,48,48]`

Train each model for 1000 epochs, learning rate 0.01, and gamma 0.995. Show the plot for each.

```

In [69]: # YOUR CODE GOES HERE
model1 = InverseArm(hidden_layer_sizes = [12], max_angle = None)
model2 = InverseArm(hidden_layer_sizes = [24, 24], max_angle = None)
model3 = InverseArm(hidden_layer_sizes = [48, 48, 48], max_angle = None)

model1 = train(model1, X_train, X_val, epochs = 1000, lr = 0.01, gamma = 0.995, create_plot = True)
model2 = train(model2, X_train, X_val, epochs = 1000, lr = 0.01, gamma = 0.995, create_plot = True)
model3 = train(model3, X_train, X_val, epochs = 1000, lr = 0.01, gamma = 0.995, create_plot = True)

```

C:\Users\barat\AppData\Local\Temp\ipykernel\_15088\664775586.py:5: UserWarning: To copy construct from a tensor, it is recommended to use `sourceTensor.clone().detach()` or `sourceTensor.clone().detach().requires_grad_(True)`, rather than `torch.tensor(sourceTensor)`.

```
X_train = torch.tensor(X_train, dtype=torch.float32)
```

C:\Users\barat\AppData\Local\Temp\ipykernel\_15088\664775586.py:6: UserWarning: To copy construct from a tensor, it is recommended to use `sourceTensor.clone().detach()` or `sourceTensor.clone().detach().requires_grad_(True)`, rather than `torch.tensor(sourceTensor)`.

```
X_val = torch.tensor(X_val, dtype=torch.float32)
```


```

Epoch    0/1000: Train Loss = 4.8425, Validation Loss = 4.4216
Epoch   40/1000: Train Loss = 0.6166, Validation Loss = 0.6107
Epoch   80/1000: Train Loss = 0.1139, Validation Loss = 0.1108
Epoch  120/1000: Train Loss = 0.0319, Validation Loss = 0.0315
Epoch  160/1000: Train Loss = 0.0223, Validation Loss = 0.0227
Epoch  200/1000: Train Loss = 0.0185, Validation Loss = 0.0187
Epoch  240/1000: Train Loss = 0.0155, Validation Loss = 0.0156
Epoch  280/1000: Train Loss = 0.0129, Validation Loss = 0.0131
Epoch  320/1000: Train Loss = 0.0108, Validation Loss = 0.0109
Epoch  360/1000: Train Loss = 0.0091, Validation Loss = 0.0092
Epoch  400/1000: Train Loss = 0.0077, Validation Loss = 0.0077
Epoch  440/1000: Train Loss = 0.0066, Validation Loss = 0.0066
Epoch  480/1000: Train Loss = 0.0056, Validation Loss = 0.0056
Epoch  520/1000: Train Loss = 0.0048, Validation Loss = 0.0048
Epoch  560/1000: Train Loss = 0.0042, Validation Loss = 0.0042
Epoch  600/1000: Train Loss = 0.0036, Validation Loss = 0.0037
Epoch  640/1000: Train Loss = 0.0032, Validation Loss = 0.0032
Epoch  680/1000: Train Loss = 0.0028, Validation Loss = 0.0028
Epoch  720/1000: Train Loss = 0.0025, Validation Loss = 0.0025
Epoch  760/1000: Train Loss = 0.0022, Validation Loss = 0.0023
Epoch  800/1000: Train Loss = 0.0020, Validation Loss = 0.0021
Epoch  840/1000: Train Loss = 0.0018, Validation Loss = 0.0019
Epoch  880/1000: Train Loss = 0.0017, Validation Loss = 0.0017
Epoch  920/1000: Train Loss = 0.0015, Validation Loss = 0.0016
Epoch  960/1000: Train Loss = 0.0014, Validation Loss = 0.0015

```




Epoch	0/1000	Train Loss = 2.6520	Validation Loss = 1.7917
Epoch	40/1000	Train Loss = 0.0806	Validation Loss = 0.0777
Epoch	80/1000	Train Loss = 0.0179	Validation Loss = 0.0185
Epoch	120/1000	Train Loss = 0.0078	Validation Loss = 0.0081
Epoch	160/1000	Train Loss = 0.0030	Validation Loss = 0.0032
Epoch	200/1000	Train Loss = 0.0017	Validation Loss = 0.0017
Epoch	240/1000	Train Loss = 0.0012	Validation Loss = 0.0013
Epoch	280/1000	Train Loss = 0.0009	Validation Loss = 0.0010
Epoch	320/1000	Train Loss = 0.0008	Validation Loss = 0.0008
Epoch	360/1000	Train Loss = 0.0007	Validation Loss = 0.0007
Epoch	400/1000	Train Loss = 0.0007	Validation Loss = 0.0007
Epoch	440/1000	Train Loss = 0.0006	Validation Loss = 0.0006
Epoch	480/1000	Train Loss = 0.0006	Validation Loss = 0.0006
Epoch	520/1000	Train Loss = 0.0005	Validation Loss = 0.0005
Epoch	560/1000	Train Loss = 0.0005	Validation Loss = 0.0005
Epoch	600/1000	Train Loss = 0.0004	Validation Loss = 0.0004
Epoch	640/1000	Train Loss = 0.0004	Validation Loss = 0.0004
Epoch	680/1000	Train Loss = 0.0004	Validation Loss = 0.0004
Epoch	720/1000	Train Loss = 0.0003	Validation Loss = 0.0003
Epoch	760/1000	Train Loss = 0.0003	Validation Loss = 0.0004
Epoch	800/1000	Train Loss = 0.0003	Validation Loss = 0.0003
Epoch	840/1000	Train Loss = 0.0003	Validation Loss = 0.0003
Epoch	880/1000	Train Loss = 0.0002	Validation Loss = 0.0002
Epoch	920/1000	Train Loss = 0.0002	Validation Loss = 0.0002
Epoch	960/1000	Train Loss = 0.0005	Validation Loss = 0.0003



```
Epoch    0/1000: Train Loss = 2.6872, Validation Loss = 2.0361
Epoch   40/1000: Train Loss = 0.0297, Validation Loss = 0.0326
Epoch   80/1000: Train Loss = 0.0021, Validation Loss = 0.0020
Epoch  120/1000: Train Loss = 0.0008, Validation Loss = 0.0008
Epoch  160/1000: Train Loss = 0.0005, Validation Loss = 0.0005
Epoch  200/1000: Train Loss = 0.0003, Validation Loss = 0.0003
Epoch  240/1000: Train Loss = 0.0002, Validation Loss = 0.0002
Epoch  280/1000: Train Loss = 0.0002, Validation Loss = 0.0002
Epoch  320/1000: Train Loss = 0.0002, Validation Loss = 0.0002
Epoch  360/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  400/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  440/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  480/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  520/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  560/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  600/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  640/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  680/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  720/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  760/1000: Train Loss = 0.0011, Validation Loss = 0.0002
Epoch  800/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  840/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  880/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  920/1000: Train Loss = 0.0001, Validation Loss = 0.0001
Epoch  960/1000: Train Loss = 0.0004, Validation Loss = 0.0004
```








## Visualizations

For each of your models, use the function `plot_predictions` to visualize model predictions on the domain. You should observe improvements with increasing network size.

In [73]: `# YOUR CODE GOES HERE`

```
plot_predictions(model1, title = "Model 1")  
plot_predictions(model2, title = "Model 2")  
plot_predictions(model3, title = "Model 3")
```





## Interactive Visualization

You can use the interactive plot below to look at the performance of your model. (The model used must be named `model`.)

```
In [72]: %matplotlib inline
from ipywidgets import interact, interactive, fixed, interact_manual, Layout, FloatSlider, Dropdown

def plot_inverse(x, y):
    xy = torch.Tensor([[x,y]])
    theta1, theta2, theta3 = model(xy).detach().numpy().flatten().tolist()
    plot_arm(theta1, theta2, theta3, show=False)
    plt.scatter(x, y, s=100, c="red", zorder=1000, marker="x")
    plt.plot([0,2,2,0],[0,0,2,2],c="lightgray",linewidth=1,zorder=-1000)
    plt.show()
```

```

slider1 = FloatSlider(value=1, min=-.5, max=2.5, step=1/100, description='x', disabled=False, continuous_update:
slider2 = FloatSlider(value=1, min=-.5, max=2.5, step=1/100, description='y', disabled=False, continuous_update:

interactive_plot = interactive(plot_inverse, x = slider1, y = slider2)
output = interactive_plot.children[-1]
output.layout.height = '600px'

interactive_plot

```

Out[72]: interactive(children=(FloatSlider(value=1.0, description='x', layout=Layout(width='550px'), max=2.5, min=-0.5,...

## Training more neural networks

Now train more networks with the following details:

1. `hidden_layer_sizes=[48,48]`, `max_angle=torch.pi/2`, train with `lr=0.01`, `epochs=1000`, `gamma=.995`
2. `hidden_layer_sizes=[48,48]`, `max_angle=None`, train with `lr=1`, `epochs=1000`, `gamma=1`
3. `hidden_layer_sizes=[48,48]`, `max_angle=2`, train with `lr=0.0001`, `epochs=300`, `gamma=1`

For each network, show a loss curve plot and a `plot_predictions` plot.

```

In [74]: # YOUR CODE GOES HERE
model4 = InverseArm(hidden_layer_sizes=[48, 48], max_angle = np.pi/2)
model4 = train(model4, X_train, X_val, epochs = 1000, lr = 0.01, gamma = 0.995, create_plot = True)
plot_predictions(model4, title = "Model 4")

model5 = InverseArm(hidden_layer_sizes = [48,48], max_angle = None)
model5 = train(model5, X_train, X_val, epochs = 1000, lr = 1, gamma = 1, create_plot = True)
plot_predictions(model5, title="Model 5")

model6 = InverseArm(hidden_layer_sizes = [48,48], max_angle = 2)
model6 = train(model6, X_train, X_val, epochs = 300, lr = 0.0001, gamma = 1, create_plot = True)
plot_predictions(model6, title = "Model 6")

```

C:\Users\barat\AppData\Local\Temp\ipykernel\_15088\664775586.py:5: UserWarning: To copy construct from a tensor, it is recommended to use `sourceTensor.clone().detach()` or `sourceTensor.clone().detach().requires_grad_(True)`, rather than `torch.tensor(sourceTensor)`.

```
X_train = torch.tensor(X_train, dtype=torch.float32)
```



C:\Users\barat\AppData\Local\Temp\ipykernel\_15088\664775586.py:6: UserWarning: To copy construct from a tensor, it is recommended to use `sourceTensor.clone().detach()` or `sourceTensor.clone().detach().requires_grad_(True)`, rather than `torch.tensor(sourceTensor)`.

```
X_val = torch.tensor(X_val, dtype=torch.float32)
```


```

Epoch    0/1000: Train Loss = 2.2229, Validation Loss = 1.5491
Epoch   40/1000: Train Loss = 0.0807, Validation Loss = 0.0795
Epoch   80/1000: Train Loss = 0.0265, Validation Loss = 0.0287
Epoch  120/1000: Train Loss = 0.0203, Validation Loss = 0.0227
Epoch  160/1000: Train Loss = 0.0190, Validation Loss = 0.0214
Epoch  200/1000: Train Loss = 0.0183, Validation Loss = 0.0207
Epoch  240/1000: Train Loss = 0.0179, Validation Loss = 0.0204
Epoch  280/1000: Train Loss = 0.0177, Validation Loss = 0.0202
Epoch  320/1000: Train Loss = 0.0175, Validation Loss = 0.0200
Epoch  360/1000: Train Loss = 0.0174, Validation Loss = 0.0198
Epoch  400/1000: Train Loss = 0.0172, Validation Loss = 0.0197
Epoch  440/1000: Train Loss = 0.0173, Validation Loss = 0.0197
Epoch  480/1000: Train Loss = 0.0171, Validation Loss = 0.0196
Epoch  520/1000: Train Loss = 0.0170, Validation Loss = 0.0195
Epoch  560/1000: Train Loss = 0.0170, Validation Loss = 0.0195
Epoch  600/1000: Train Loss = 0.0169, Validation Loss = 0.0194
Epoch  640/1000: Train Loss = 0.0170, Validation Loss = 0.0195
Epoch  680/1000: Train Loss = 0.0169, Validation Loss = 0.0193
Epoch  720/1000: Train Loss = 0.0169, Validation Loss = 0.0193
Epoch  760/1000: Train Loss = 0.0169, Validation Loss = 0.0193
Epoch  800/1000: Train Loss = 0.0169, Validation Loss = 0.0194
Epoch  840/1000: Train Loss = 0.0168, Validation Loss = 0.0192
Epoch  880/1000: Train Loss = 0.0168, Validation Loss = 0.0192
Epoch  920/1000: Train Loss = 0.0170, Validation Loss = 0.0193
Epoch  960/1000: Train Loss = 0.0167, Validation Loss = 0.0192

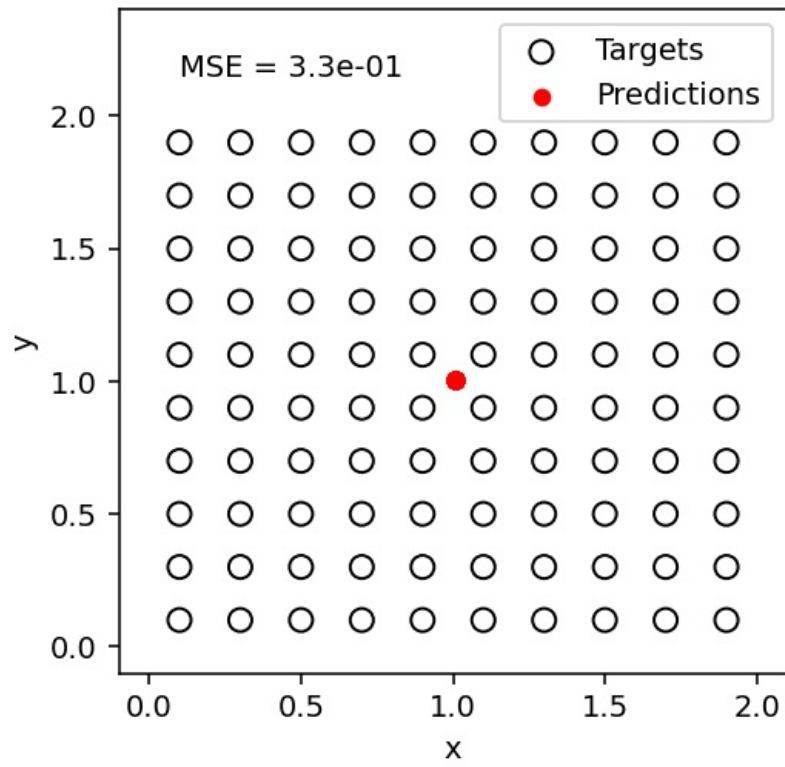
```




Epoch 0/1000: Train Loss = 2.6804, Validation Loss = 2.8504  
Epoch 40/1000: Train Loss = 0.3415, Validation Loss = 0.8446  
Epoch 80/1000: Train Loss = 0.5532, Validation Loss = 0.4691  
Epoch 120/1000: Train Loss = 0.3397, Validation Loss = 0.3425  
Epoch 160/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 200/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 240/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 280/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 320/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 360/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 400/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 440/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 480/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 520/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 560/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 600/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 640/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 680/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 720/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 760/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 800/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 840/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 880/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 920/1000: Train Loss = 0.3395, Validation Loss = 0.3426  
Epoch 960/1000: Train Loss = 0.3395, Validation Loss = 0.3426




Model 5



Epoch 0/300: Train Loss = 2.4209, Validation Loss = 2.4439  
Epoch 12/300: Train Loss = 2.1337, Validation Loss = 2.1677  
Epoch 24/300: Train Loss = 1.9155, Validation Loss = 1.9580  
Epoch 36/300: Train Loss = 1.7620, Validation Loss = 1.8101  
Epoch 48/300: Train Loss = 1.6569, Validation Loss = 1.7079  
Epoch 60/300: Train Loss = 1.5812, Validation Loss = 1.6333  
Epoch 72/300: Train Loss = 1.5192, Validation Loss = 1.5718  
Epoch 84/300: Train Loss = 1.4634, Validation Loss = 1.5162  
Epoch 96/300: Train Loss = 1.4113, Validation Loss = 1.4644  
Epoch 108/300: Train Loss = 1.3624, Validation Loss = 1.4158  
Epoch 120/300: Train Loss = 1.3157, Validation Loss = 1.3692  
Epoch 132/300: Train Loss = 1.2704, Validation Loss = 1.3239  
Epoch 144/300: Train Loss = 1.2264, Validation Loss = 1.2800  
Epoch 156/300: Train Loss = 1.1829, Validation Loss = 1.2364  
Epoch 168/300: Train Loss = 1.1385, Validation Loss = 1.1918  
Epoch 180/300: Train Loss = 1.0926, Validation Loss = 1.1456  
Epoch 192/300: Train Loss = 1.0450, Validation Loss = 1.0977  
Epoch 204/300: Train Loss = 0.9961, Validation Loss = 1.0483  
Epoch 216/300: Train Loss = 0.9467, Validation Loss = 0.9982  
Epoch 228/300: Train Loss = 0.8975, Validation Loss = 0.9483  
  
Epoch 240/300: Train Loss = 0.8496, Validation Loss = 0.8995  
Epoch 252/300: Train Loss = 0.8037, Validation Loss = 0.8526  
Epoch 264/300: Train Loss = 0.7604, Validation Loss = 0.8082  
Epoch 276/300: Train Loss = 0.7202, Validation Loss = 0.7667  
Epoch 288/300: Train Loss = 0.6831, Validation Loss = 0.7282





## Prompts

None of these 3 models should have great performance. Describe what went wrong in each case.

Model 4 is performing poorly because the `max_angle` parameter is incorrectly set. Since `max_angle` is set to  $\pi/2$ , the predictions for the lower half of the x and y coordinates are incorrect.

Model 5 is underperforming due to an excessively high learning rate of 1. This causes the model to oscillate around the global minimum without converging.

Model 6 is underfitting because the number of epochs is too low for the model to reach the global optimum. As a result, the loss function does not fully converge, leading to poor performance.